

Homework 3

Part 1: Imbalanced Dataset

- In this homework, you will be working with an imbalanced Dataset.
- The dataset is Credit Card Fraud Detection dataset which was hosted on Kaggle.
- The aim is to detect fraudulent transactions.

Instructions

- 1) Please push the .ipynb and .pdf to Github Classroom prior to the deadline, .py file is optional (not needed).
- 2) Please include your Name and UNI below.

Name: Shruti Agarwal

UNI: sa4136

Setup

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: # Feel free to import any other packages you need
from sklearn.model_selection import train_test_split, cross_validate
from sklearn.preprocessing import StandardScaler
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import cross_val_score
from sklearn.metrics import roc_auc_score, average_precision_score

# !pip install imbalanced-learn
from imblearn.over_sampling import RandomOverSampler
from imblearn.under_sampling import RandomUnderSampler
from imblearn.over_sampling import SMOTE
from sklearn.metrics import confusion_matrix, plot_confusion_matrix
from sklearn.metrics import roc_curve
```

Data Preprocessing and Exploration.

- Download the Kaggle Credit Card Fraud data set.
- Features V1, V2, ... V27, V28 are the principal components obtained with PCA, the only features which have not been transformed with PCA are 'Time' and 'Amount'.
- Feature 'Time' contains the seconds elapsed between each transaction and the first transaction in the dataset.
- The feature 'Amount' is the transaction Amount, this feature can be used for example-dependant cost-sensitive learning.
- Feature 'Class' is the response variable and it takes value 1 in case of fraud and 0 otherwise.

```
In [3]: raw_df = pd.read_csv('https://storage.googleapis.com/download.tensorflow.org/sample-credit-card-fraud/sample_000001.csv')
raw_df.head(10)
```

Out [3]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533
5	2.0	-0.425966	0.960523	1.141109	-0.168252	0.420987	-0.029728	0.476201	0.260314
6	4.0	1.229658	0.141004	0.045371	1.202613	0.191881	0.272708	-0.005159	0.081213
7	7.0	-0.644269	1.417964	1.074380	-0.492199	0.948934	0.428118	1.120631	-3.807864
8	7.0	-0.894286	0.286157	-0.113192	-0.271526	2.669599	3.721818	0.370145	0.851084
9	9.0	-0.338262	1.119593	1.044367	-0.222187	0.499361	-0.246761	0.651583	0.069539

10 rows × 31 columns

Examining the class Imbalance

1.1 How many observations are in this dataset? How many are positive and negative?

(Note: Positive labels are labeled as 1)

In [4]: *# Your Code Here*

```
num_obs = raw_df.shape[0]
obs = raw_df.groupby(['Class']).size()
pos_obs = obs[1]
neg_obs = obs[0]
print(f"There are {num_obs} observations in the dataset with {pos_obs}
```

There are 284807 observations in the dataset with 492 positive and 284315 negative observations.

1.2 Cleaning and normalizing the data

The raw data has a few issues.

Since we are unsure what the time column actually means so drop the Time column. The Amount column also has a wide range of values covered so we take the log of the Amount column to reduce its range.

The below is already done for you.

In [5]: `cleaned_df = raw_df.copy()`

```
# You don't want the 'Time' column. Pop it off
cleaned_df.pop('Time')

# The 'Amount' column covers a huge range. Convert it to log-space.
eps = 0.001
cleaned_df['Log Ammount'] = np.log(cleaned_df.pop('Amount') + eps)
```

1.2.1 Split the dataset into development and test sets. Set test size as 20% and random state as 42. Print the shape of your development and test features

In [6]: *# Your Code Here*

```
X = cleaned_df.drop('Class', axis=1)
y = cleaned_df['Class']

# split the dataset into development and test sets
X_dev, X_test, y_dev, y_test = train_test_split(X, y, test_size=0.2, r

# print the shape of the development and test sets
print("Development set features shape:", X_dev.shape)
print("Test set features shape:", X_test.shape)
```

Development set features shape: (227845, 29)

Test set features shape: (56962, 29)

1.2.2 Normalize the features using Standard Scaler from Sklearn.

In [7]: *# Your Code Here*

```
scaler = StandardScaler()
X_dev_scaled = scaler.fit_transform(X_dev)
X_test_scaled = scaler.transform(X_test)
```

Default Baseline

1.3.1 First, let us fit a default Decision tree classifier (use max_depth=10 and random_state=42). Print the AUC and Average Precision values of 5 Fold Cross Validation

In [8]: *# Your Code Here*

```
dt_def = DecisionTreeClassifier(max_depth=10, random_state=42)

auc_scores = cross_val_score(dt_def, X_dev_scaled, y_dev, cv=5, scoring='roc_auc')
ap_scores = cross_val_score(dt_def, X_dev_scaled, y_dev, cv=5, scoring='average_precision')

# print the AUC and Average Precision scores
print("AUC scores:", auc_scores)
print("Average Precision scores:", ap_scores)
print("Average AUC:", auc_scores.mean())
print("Average Average Precision:", ap_scores.mean())
```

```
AUC scores: [0.86957634 0.82799822 0.83798122 0.88266235 0.90893353]
Average Precision scores: [0.64646286 0.6763446 0.61677062 0.68317716 0.68281098]
Average AUC: 0.8654303324923142
Average Average Precision: 0.6611132453729376
```

Random Oversampling

1.3.2 Perform random oversampling on the development dataset.

- How many positive and negative labels do you observe after random oversampling?
- What is the shape of your development dataset?

(Note: Set random state as 42 when performing oversampling)

In [9]: *# Your Code Here*

```
ros = RandomOverSampler(random_state=42)
X_dev_ros, y_dev_ros = ros.fit_resample(X_dev_scaled, y_dev)

print("Number of positive labels:", sum(y_dev_ros == 1))
print("Number of negative labels:", sum(y_dev_ros == 0))
print("Shape of the resampled development set:", X_dev_ros.shape)
```

```
Number of positive labels: 227451
Number of negative labels: 227451
Shape of the resampled development set: (454902, 29)
```

1.3.3 Repeat 1.3.1 using the dataset you created in the above step (1.3.2 Random oversampling).

(Make sure you use the same hyperparameters as 1.3.1. i.e., max_depth=10, random_state=42 and 5 Fold Cross Validation)

This will help us to compare the models.

In [10]: *# Your Code Here*

```
dt_ros = DecisionTreeClassifier(max_depth=10, random_state=42)

auc_scores = cross_val_score(dt_ros, X_dev_ros, y_dev_ros, cv=5, scoring='roc_auc')
ap_scores = cross_val_score(dt_ros, X_dev_ros, y_dev_ros, cv=5, scoring='average_precision')

# print the AUC and Average Precision scores
print("AUC scores:", auc_scores)
print("Average Precision scores:", ap_scores)
print("Average AUC:", auc_scores.mean())
print("Average Average Precision:", ap_scores.mean())
```

```
AUC scores: [0.99840012 0.9983086  0.99847279 0.99847804 0.99857557]
Average Precision scores: [0.99786268 0.99767719 0.9978636  0.9978744
4 0.99807898]
Average AUC: 0.9984470244132799
Average Average Precision: 0.9978713776423949
```

Random Undersampling

1.3.4 Perform Random undersampling on the development dataset.

- How many positive and negative labels do you observe after random undersampling?
- What is the shape of your development dataset?

(Note: Set random state as 42 when performing undersampling)

In [11]: *# Your Code Here*

```
rus = RandomUnderSampler(random_state=42)
X_dev_rus, y_dev_rus = rus.fit_resample(X_dev_scaled, y_dev)
```

In [12]: `print("Number of positive labels:", sum(y_dev_rus == 1))`
`print("Number of negative labels:", sum(y_dev_rus == 0))`

Number of positive labels: 394
Number of negative labels: 394

In [13]: `print("Shape of the resampled development set:", X_dev_rus.shape)`

Shape of the resampled development set: (788, 29)

1.3.5 Repeat 1.3.1 using the dataset you created in the above step(1.3.4 Random undersampling).

(Make sure you use the same hyperparameters as 1.3.1. i.e., max_depth=10, random_state=42 and 5 Fold Cross Validation)

This will help us to compare the models

In [14]: *# Your Code Here*

```
dt_rus = DecisionTreeClassifier(max_depth=10, random_state=42)

auc_scores = cross_val_score(dt_rus, X_dev_rus, y_dev_rus, cv=5, scoring='roc_auc')
ap_scores = cross_val_score(dt_rus, X_dev_rus, y_dev_rus, cv=5, scoring='average_precision')

# print the AUC and Average Precision scores
print("AUC scores:", auc_scores)
print("Average Precision scores:", ap_scores)
print("Average AUC:", auc_scores.mean())
print("Average Average Precision:", ap_scores.mean())
```

AUC scores: [0.8852748 0.86388399 0.89873418 0.91731581 0.94790652]
Average Precision scores: [0.82838644 0.81661392 0.86683544 0.87531025 0.92073059]
Average AUC: 0.9026230592566116
Average Average Precision: 0.8615753308133469

SMOTE

1.3.6 Perform Synthetic Minority Oversampling Technique (SMOTE) on the development dataset

- How many positive and negative labels do you observe after performing SMOTE?
- What is the shape of your development dataset?

(Note: Set random state as 42 when performing SMOTE)

In [15]: *# Your Code Here*

```
smote = SMOTE(random_state=42)
X_dev_smote, y_dev_smote = smote.fit_resample(X_dev_scaled, y_dev)
```

In [16]:

```
print("Number of positive labels:", sum(y_dev_smote == 1))
print("Number of negative labels:", sum(y_dev_smote == 0))
```

```
Number of positive labels: 227451
Number of negative labels: 227451
```

In [17]:

```
print("Shape of the resampled development set:", X_dev_smote.shape)
```

```
Shape of the resampled development set: (454902, 29)
```

1.3.7 Repeat 1.3.1 using the dataset you created in the above step(1.3.4 SMOTE).

(Make sure you use the same hyperparameters as 1.3.1. i.e., max_depth=10, random_state=42 and 5 Fold Cross Validation)

This will help us to compare the models

In [18]: *# Your Code Here*

```
dt_smote = DecisionTreeClassifier(max_depth=10, random_state=42)

auc_scores = cross_val_score(dt_smote, X_dev_smote, y_dev_smote, cv=5,
                               ap_scores = cross_val_score(dt_smote, X_dev_smote, y_dev_smote, cv=5,

# print the AUC and Average Precision scores
print("AUC scores:", auc_scores)
print("Average Precision scores:", ap_scores)
print("Average AUC:", auc_scores.mean())
print("Average Average Precision:", ap_scores.mean())
```

```
AUC scores: [0.99729603 0.99747394 0.99713467 0.99703584 0.99724342]
Average Precision scores: [0.99658165 0.99656764 0.99603047 0.9959331
7 0.99618909]
Average AUC: 0.9972367791948902
Average Average Precision: 0.996260402041384
```

Balanced Weight

1.3.8 Train a balanced default Decision tree classifier.

[use max_depth=10 and random_state=42 and balance the class weights with 5 Fold Cross Validation]

Print the AUC and average precision on dev set

In [19]: *# Your Code Here*

```
dt_balanced = DecisionTreeClassifier(max_depth=10, random_state=42, cl
cv_results = cross_validate(dt_balanced, X_dev_scaled, y_dev, cv=5, sc

# print the AUC and average precision on the development set
print("AUC on development set:", cv_results['test_roc_auc'].mean())
print("Average precision on development set:", cv_results['test_averag
```

AUC on development set: 0.9070005244195443

Average precision on development set: 0.623661060432623

Model Prediction & Evaluation

1.4.1 Make predictions on the test set using the five models that you built and report their AUC values.

(Five models include models from - Default Baseline, Random Undersampling, Random Oversampling, SMOTE & Balanced Weight)

In [20]: *# Your Code Here*

```
dt_def.fit(X_dev_scaled, y_dev)
dt_ros.fit(X_dev_ros, y_dev_ros)
dt_rus.fit(X_dev_rus, y_dev_rus)
dt_smote.fit(X_dev_smote, y_dev_smote)
dt_balanced.fit(X_dev_scaled, y_dev)

# make predictions on the test set and compute AUC scores
y_pred_def = dt_def.predict_proba(X_test_scaled)[: , 1]
auc_def = roc_auc_score(y_test, y_pred_def)

y_pred_ros = dt_ros.predict_proba(X_test_scaled)[: , 1]
auc_ros = roc_auc_score(y_test, y_pred_ros)

y_pred_rus = dt_rus.predict_proba(X_test_scaled)[: , 1]
auc_rus = roc_auc_score(y_test, y_pred_rus)

y_pred_smote = dt_smote.predict_proba(X_test_scaled)[: , 1]
auc_smote = roc_auc_score(y_test, y_pred_smote)

y_pred_balanced = dt_balanced.predict_proba(X_test_scaled)[: , 1]
auc_balanced = roc_auc_score(y_test, y_pred_balanced)

# print the AUC scores
print("AUC for Default Baseline:", auc_def)
print("AUC for Random Oversampling:", auc_ros)
print("AUC for Random Undersampling:", auc_rus)
print("AUC for Smote:", auc_smote)
print("AUC for Balanced Weights:", auc_balanced)
```

```
AUC for Default Baseline: 0.871912253224306
AUC for Random Oversampling: 0.8609190528349776
AUC for Random Undersampling: 0.9050787485787788
AUC for Smote: 0.8893473364303515
AUC for Balanced Weights: 0.8432507780827582
```

1.4.2 Plot Confusion Matrices for all the five models on the test set. Comment your results and share your observations of the confusion matrices in detail

(Five models include models from - Default Baseline, Random Undersampling, Random Oversampling, SMOTE & Balanced Weight)

In [21]:

```

titles = ['Default Baseline', 'Random Oversampling', 'Random Undersamp
classifiers = [dt_def, dt_ros, dt_rus, dt_smote, dt_balanced]

```

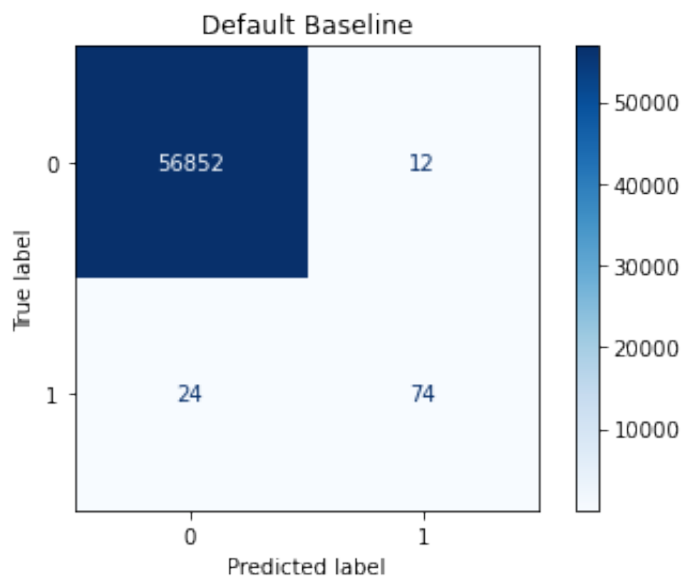
```

for clf, title in zip(classifiers, titles):
    # predict the target variable for the test set
    y_pred = clf.predict(X_test_scaled)
    # calculate the confusion matrix
    cm = confusion_matrix(y_test, y_pred)
    # plot the confusion matrix
    plot_confusion_matrix(clf, X_test_scaled, y_test, cmap=plt.cm.Blue
plt.title(title)
plt.show()

```


/Users/shrutiagarwal/opt/anaconda3/lib/python3.9/site-packages/sklearn/
utils/deprecation.py:87: FutureWarning: Function plot_confusion_mat
rix is deprecated; Function `plot_confusion_matrix` is deprecated in
1.0 and will be removed in 1.2. Use one of the class methods: Confusi
onMatrixDisplay.from_predictions or ConfusionMatrixDisplay.from_estim
ator.


```
warnings.warn(msg, category=FutureWarning)
```



/Users/shrutiagarwal/opt/anaconda3/lib/python3.9/site-packages/sklearn/
utils/deprecation.py:87: FutureWarning: Function plot_confusion_mat
rix is deprecated; Function `plot_confusion_matrix` is deprecated in
1.0 and will be removed in 1.2. Use one of the class methods: Confusi
onMatrixDisplay.from_predictions or ConfusionMatrixDisplay.from_estim
ator.

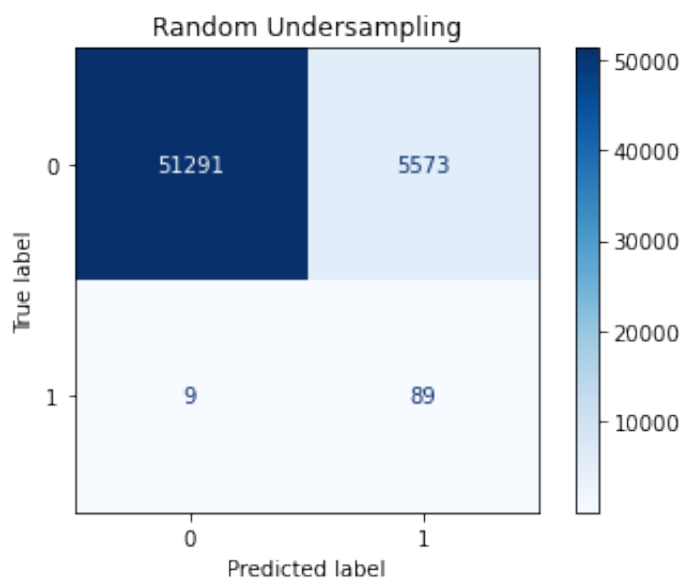
```
warnings.warn(msg, category=FutureWarning)
```






/Users/shrutiagarwal/opt/anaconda3/lib/python3.9/site-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function plot_confusion_matrix is deprecated; Function `plot_confusion_matrix` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: ConfusionMatrixDisplay.from_predictions or ConfusionMatrixDisplay.from_estimator.


warnings.warn(msg, category=FutureWarning)



/Users/shrutiagarwal/opt/anaconda3/lib/python3.9/site-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function plot_confusion_matrix is deprecated; Function `plot_confusion_matrix` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: ConfusionMatrixDisplay.from_predictions or ConfusionMatrixDisplay.from_estimator.


warnings.warn(msg, category=FutureWarning)





/Users/shrutiagarwal/opt/anaconda3/lib/python3.9/site-packages/sklearn/Utils/deprecation.py:87: FutureWarning: Function plot_confusion_matrix is deprecated; Function `plot_confusion_matrix` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: ConfusionMatrixDisplay.from_predictions or ConfusionMatrixDisplay.from_estimator.

warnings.warn(msg, category=FutureWarning)

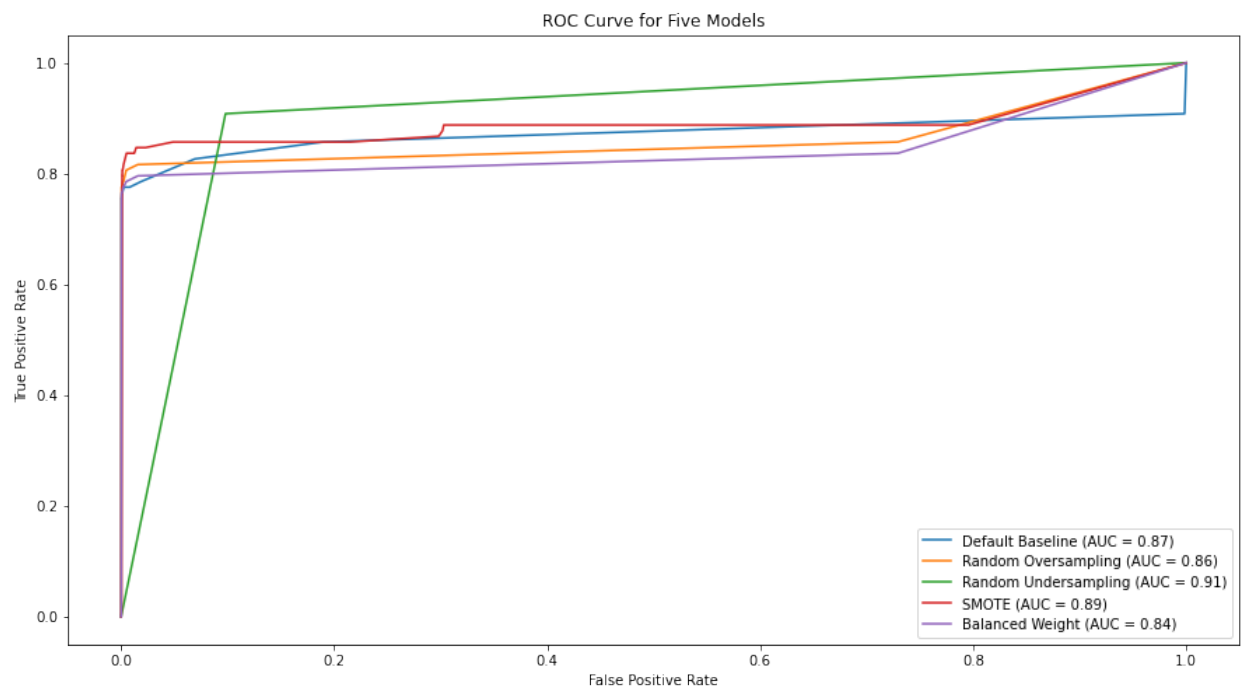


1.4.3 Plot ROC for all the five models on the test set in a single plot. Make sure you label axes and legend properly. Comment on your results and share your observations in detail

(Five models include models from - Default Baseline, Random Undersampling, Random Oversampling, SMOTE & Balanced Weight)

In [22]: *# Your Code Here*

```
titles = ['Default Baseline', 'Random Oversampling', 'Random Undersamp  
classifiers = [dt_def, dt_ros, dt_rus, dt_smote, dt_balanced]  
  
plt.figure(figsize=(15, 8))  
for clf, title in zip(classifiers, titles):  
    y_prob = clf.predict_proba(X_test_scaled)[: , 1]  
    fpr, tpr, thresholds = roc_curve(y_test, y_prob)  
    auc = roc_auc_score(y_test, y_prob)  
    plt.plot(fpr, tpr, label=f'{title} (AUC = {auc:.2f})')  
  
# add labels and legend  
plt.xlabel('False Positive Rate')  
plt.ylabel('True Positive Rate')  
plt.title('ROC Curve for Five Models')  
plt.legend()  
plt.show()
```



In []: